

The side predicates

How `side1` to `side4` are derived, and how they are evaluated

4 October 2026

What this is. A reference to come back to when the predicates of `kernel/side.h` stop making sense. It derives every formula from scratch, checks each one on a worked example (the same examples as the hand-written tests), and says how the formula is turned into an exact, filtered predicate. It contains no code, on purpose. Every formula here was also checked against direct rational construction of the vertex on random inputs.

Contents

1	What a side predicate answers	2
1.1	The question	2
1.2	The rule: never construct the vertex	2
2	The building block: f_{ia}	3
2.1	It is affine	3
2.2	The form to evaluate it in	3
2.3	Degrees	3
3	The one lemma behind <code>side1</code> to <code>side3</code>	4
4	<code>side1</code>: a domain corner	5
5	<code>side2</code>: an edge crossing one bisector	5
5.1	Derivation	5
5.2	The denominator, explained	6
5.3	Worked example	6
5.4	Symmetries the tests check	6
6	<code>side3</code>: a triangle crossing two bisectors	6
6.1	Derivation	7
6.2	Structure: share the bottom rows	7
6.3	Worked example	8
6.4	Symmetries the tests check	8
6.5	Cost of the exact path	8
7	<code>side4</code>: a Voronoi vertex	8
7.1	Work relative to p_i	9
7.2	Derivation	9
7.3	Structure: four triple products	9
7.4	The <code>orient3d</code> trap	10
7.5	Worked example	10
7.6	Symmetries the tests check	10

8	Evaluating them: fast when easy, exact when not	10
8.1	One formula, several number types	10
8.2	The filter	11
8.3	The decision	11
8.4	Why <code>DynamicExpansion</code>	11
8.5	Counters	12
9	How they are tested	12
9.1	The oracle	12
9.2	The input families	12
9.3	Thresholds	12
10	A preview: ties	13
11	Checklist	13

1 What a side predicate answers

1.1 The question

The clipper holds a convex piece of seed i 's cell and wants to cut it by the bisector of i and a new seed m . For each vertex x of the piece it asks one question:

Is x on i 's side of the bisector of i and m , on m 's side, or exactly on it?

With the power distance $\pi_a(x) = \|x - p_a\|^2 - w_a$, the answer is the sign of

$$f_{im}(x) = \pi_i(x) - \pi_m(x) :$$

- 1 x is closer to i (in power distance): it stays
- +1 x is closer to m : it is cut away
- 0 x is exactly on the bisector: a tie (§10)

1.2 The rule: never construct the vertex

The obvious approach computes the coordinates of x and then evaluates $f_{im}(x)$. That is what a floating-point Voronoi code does, and it is exactly what this design forbids (DESIGN.tex, section 6.1). A constructed x is rounded, and two cells that look at the same vertex through different roundings can take different decisions. That is how cracks, asymmetric neighbours and broken dual meshes happen.

Instead, each predicate is the sign of a *polynomial in the input numbers*: seed coordinates, weights and domain-vertex coordinates. The sign of a polynomial in doubles can be computed exactly. The vertex is never constructed. Only its *definition* is used:

type	vertex defined by	lives on	predicate
V	a domain corner q	a domain vertex	side1
E	domain edge q_0q_1 and bisector (i, j)	a domain edge	side2
F	domain triangle $q_0q_1q_2$ and bisectors $(i, j), (i, k)$	a domain face	side3
C	bisectors $(i, j), (i, k), (i, l)$	nowhere special: a Voronoi vertex	side4

Notice the pattern: a vertex on a feature with k free directions (0 for a corner, 1 for an edge, 2 for a face, 3 for open space) needs k bisector equations to be pinned down.

2 The building block: f_{ia}

2.1 It is affine

Expand the two power distances:

$$f_{ia}(x) = \|x\|^2 - 2x \cdot p_i + \|p_i\|^2 - w_i - \|x\|^2 + 2x \cdot p_a - \|p_a\|^2 + w_a.$$

The $\|x\|^2$ terms cancel, which leaves

$$f_{ia}(x) = 2x \cdot (p_a - p_i) - c_a, \quad c_a = \|p_a\|^2 - \|p_i\|^2 + w_i - w_a. \quad (1)$$

This is the most important fact in this document. f_{ia} **is an affine function of x** : a linear part plus a constant. When $p_a \neq p_i$ its zero set is a plane, the bisector; it is negative on i 's side and positive on a 's. When the two seeds sit at the same position, the linear part vanishes and f_{ia} is the constant $w_a - w_i$: its zero set is empty, or all of space if the weights are equal too. Such a pair has no bisector, and every vertex defined by one is singular (below).

Affine functions commute with affine combinations. If $x = \sum_r \lambda_r q_r$ with $\sum_r \lambda_r = 1$, then

$$f_{ia}(x) = \sum_r \lambda_r f_{ia}(q_r). \quad (2)$$

The condition $\sum_r \lambda_r = 1$ is what lets the constant c_a come out right: $\sum_r \lambda_r c_a = c_a$. Everything in `side1` to `side3` follows from (2).

2.2 The form to evaluate it in

Equation (1) is good for thinking and bad for computing. Its terms $\|p_a\|^2$ and $\|p_i\|^2$ are huge when the points are far from the origin, and they cancel. A rounded evaluation then loses everything, and the filter's error bound is as large as $\|p\|^2$ rather than as large as the configuration. The same quantity can be written with differences only:

$$\pi_i(q) - \pi_a(q) = \|q - p_i\|^2 - \|q - p_a\|^2 + (w_a - w_i) = (p_a - p_i) \cdot ((q - p_i) + (q - p_a)) + (w_a - w_i), \quad (3)$$

using $\|u\|^2 - \|v\|^2 = (u - v) \cdot (u + v)$ with $u = q - p_i$ and $v = q - p_a$, so that $u - v = p_a - p_i$. This is the form of `bisector_value` in `side.h`. Translating every input by the same vector changes none of the differences in exact arithmetic. If the translated coordinates are themselves exact doubles, the computed differences are the same doubles too, so the value and its error bound depend only on the *shape* of the configuration, not on where it sits. (A translation that rounds the coordinates changes the input, and with it, possibly, the answer.) The tests use exactly representable translations: they move the small-integer lattice inputs by 2^{20} and by $10^6 + 0.5$, and check that the *signs* do not change.

2.3 Degrees

Scale all coordinates by s and all weights by s^2 . Power distances scale by s^2 , so every f value scales by s^2 . So f is **degree 2** in the inputs, with weights counting as degree 2, like a squared coordinate. Degrees matter for two practical reasons:

- the higher the degree, the more rounding error the filter's bound accumulates, so the filter decides a little less often (`side3`'s threshold is the lowest);
- the higher the degree, the longer the exact expansions get, which is why the exact path uses `DynamicExpansion` and not `Expansion<N>` (§8).

3 The one lemma behind `side1` to `side3`

Lemma 1 (Cramer for affine functions). *Let $q_0, \dots, q_k$ be the corners of a domain feature ($k = 0, 1, 2$). Let $h_1, \dots, h_k$ and g be affine functions. Let x be the point of the feature's affine hull where $h_1(x) = \dots = h_k(x) = 0$. Write $h_s^{(r)} = h_s(q_r)$ and $g^{(r)} = g(q_r)$, and let*

$$M = \begin{pmatrix} 1 & 1 & \dots & 1 \\ h_1^{(0)} & h_1^{(1)} & \dots & h_1^{(k)} \\ \vdots & & & \vdots \\ h_k^{(0)} & h_k^{(1)} & \dots & h_k^{(k)} \end{pmatrix}.$$

If $\det M \neq 0$, then x exists and is unique, and

$$g(x) = \frac{\det M_g}{\det M},$$

where M_g is M with its first row (the row of ones) replaced by $(g^{(0)}, \dots, g^{(k)})$.

Proof. Write x in barycentric coordinates, $x = \sum_r \lambda_r q_r$ with $\sum_r \lambda_r = 1$. By (2), $h_s(x) = \sum_r \lambda_r h_s^{(r)}$. So the conditions on x are

$$\sum_r \lambda_r = 1, \quad \sum_r \lambda_r h_s^{(r)} = 0 \quad (s = 1..k),$$

which is the square linear system $M\lambda = e_1$, where $e_1 = (1, 0, \dots, 0)^\top$. If $\det M \neq 0$ it has exactly one solution. By Cramer's rule,

$$\lambda_r = \frac{\det M^{[r]}}{\det M},$$

where $M^{[r]}$ is M with column r replaced by e_1 . Expanding $\det M^{[r]}$ along that column leaves only its top entry, so $\det M^{[r]} = C_{0r}$, the cofactor of the entry in row 0, column r of M . Then, by (2) again,

$$g(x) = \sum_r \lambda_r g^{(r)} = \frac{\sum_r g^{(r)} C_{0r}}{\det M}.$$

The numerator is the expansion along row 0 of a determinant whose row 0 is $(g^{(0)}, \dots, g^{(k)})$ and whose other rows are those of M , which is $\det M_g$. The cofactors C_{0r} do not involve row 0, so changing it does not change them. $\square$

What it means for the predicates. Take h_s to be the bisector functions that define the vertex $(f_{ij}$, and f_{ik} for a face) and $g = f_{im}$. Then

$$\text{sign } f_{im}(x) = \text{sign}(\det M_g) \cdot \text{sign}(\det M).$$

Both determinants are polynomials in f values at the corners, so they are polynomials in the input. That is the whole trick. Three remarks:

- **The denominator's sign is not fixed.** Listing the corners in another order permutes the columns of M and M_g together. Both determinants can flip sign, and their ratio cannot. So the predicate must multiply the two signs: it cannot assume $\det M > 0$.
- **$\det M = 0$ means the equations do not define a unique vertex:** they have no solution, or infinitely many. No solution when the feature is parallel to a bisector it does not lie in; infinitely many when it lies in that bisector, or the bisectors meet it in a line instead of a point; and the feature itself may be degenerate (coincident or collinear corners), or two seeds coincide (§2). The clipper never asks about such a vertex, because it never creates one. The oracle detects exactly the same cases: its linear system has determinant $\det M$, as §5.2 shows for `side2`.

- **x does not have to lie inside the feature.** The lemma works on the whole line or plane through it, and some λ_r may be negative. The predicate does not care. The tests use random features whose vertex is often outside them.

4 side1: a domain corner

The case $k = 0$ of the lemma: $M = (1)$ and $M_g = (g^{(0)})$, so

$$f_{im}(x) = f_{im}(q).$$

The vertex is the corner q itself, nothing needs solving, and the predicate is the sign of one f value, in the difference form (3). Degree 2.

Worked example (from `test_side1.cpp`). $p_i = 0$, $p_m = (2, 0, 0)$, no weights, $q = (\frac{1}{2}, 0, 0)$. Then $p_m - p_i = (2, 0, 0)$ and $(q - p_i) + (q - p_m) = (\frac{1}{2}, 0, 0) + (-\frac{3}{2}, 0, 0) = (-1, 0, 0)$, so $f_{im}(q) = -2$: answer -1 . Directly: $\pi_i(q) = \frac{1}{4}$ and $\pi_m(q) = \frac{9}{4}$, a difference of -2 .

Symmetry. $f_{mi} = -f_{im}$, so swapping i and m flips the sign. The test checks this on 50 000 inputs.

5 side2: an edge crossing one bisector




Figure 1: An E vertex: where the domain edge q_0q_1 crosses the bisector of i and j . **side2** asks on which side of a *third* bisector, (i, m) , it lies.

Signature (fixed by the test): `side2(pi, wi, pj, wj, pm, wm, q0, q1)`.

5.1 Derivation

Name the four f values at the corners:

$$\begin{array}{lll} a_0 = f_{ij}(q_0), & a_1 = f_{ij}(q_1) & \text{(the bisector that defines } x), \\ b_0 = f_{im}(q_0), & b_1 = f_{im}(q_1) & \text{(the bisector tested).} \end{array}$$

Directly. Write $x = (1 - t)q_0 + tq_1$. Then $f_{ij}(x) = (1 - t)a_0 + ta_1 = 0$ gives $t = a_0/(a_0 - a_1)$, and

$$f_{im}(x) = (1 - t)b_0 + tb_1 = \frac{b_0(a_0 - a_1) - a_0b_0 + a_0b_1}{a_0 - a_1} = \frac{a_0b_1 - a_1b_0}{a_0 - a_1}.$$

By the lemma, with $k = 1$, $h_1 = f_{ij}$, $g = f_{im}$:

$$f_{im}(x) = \frac{\begin{vmatrix} b_0 & b_1 \\ a_0 & a_1 \end{vmatrix}}{\begin{vmatrix} 1 & 1 \\ a_0 & a_1 \end{vmatrix}} = \frac{b_0 a_1 - b_1 a_0}{a_1 - a_0}. \quad (4)$$

It is the same expression: multiply the top and bottom of the direct one by -1 .

$\text{sign } f_{im}(x) = \text{sign}(N) \cdot \text{sign}(D), \quad N = b_0 a_1 - b_1 a_0 \text{ (degree 4)}, \quad D = a_1 - a_0 \text{ (degree 2)}.$

Cost: four f evaluations, one 2×2 determinant, one difference.

5.2 The denominator, explained

By (1), $a_1 - a_0 = f_{ij}(q_1) - f_{ij}(q_0) = 2(q_1 - q_0) \cdot (p_j - p_i)$: the constants cancel. So $D = 0$ exactly when the edge is perpendicular to $p_j - p_i$, which means parallel to the bisector (or $p_j = p_i$). Then there is no single crossing: none if the edge's line misses the bisector, a whole line of them if it lies inside it. This is also exactly the single coefficient of the oracle's 1×1 system: the same quantity, written differently.

This also gives an alternative way to compute D , straight from coordinates. Its value is identical. In **Bounded** its error bound is smaller, because it is one dot product of differences rather than a difference of two degree-2 results. That is optional and only worth it if the filter rate disappoints. If you use it, keep its sign consistent with N 's. Both forms above are $a_1 - a_0$, not $a_0 - a_1$.

5.3 Worked example

The same numbers as the hand-written test in `test_side2.cpp`.

$p_i = 0$, $p_j = (2, 0, 0)$, edge $q_0 = (-1, 0, 0)$ to $q_1 = (3, 0, 0)$, $p_m = (5, 0, 0)$, no weights. With $p_i = 0$, $f_{ia}(q) = 2q \cdot p_a - \|p_a\|^2$:

$$\begin{array}{lll} f_{ij}(q) = 4q_x - 4 : & a_0 = -8, & a_1 = 8, \\ f_{im}(q) = 10q_x - 25 : & b_0 = -35, & b_1 = 5. \end{array}$$

$N = (-35)(8) - (5)(-8) = -240$ and $D = 8 - (-8) = 16$, so $f_{im}(x) = -15$: answer -1 .

Check: the crossing is $x = (1, 0, 0)$, with $\pi_i(x) = 1$ and $\pi_m(x) = 16$, a difference of -15 . The direct alternative for D gives $2(4, 0, 0) \cdot (2, 0, 0) = 16$. ✓

5.4 Symmetries the tests check

- **The edge read backwards.** Swapping q_0 and q_1 swaps $a_0 \leftrightarrow a_1$ and $b_0 \leftrightarrow b_1$. That negates N and negates D , so the ratio is unchanged. This is the reason for multiplying the two signs.
- **i and j exchanged.** The vertex is the same point, where $f_{ji} = -f_{ij} = 0$. And $f_{jm} = \pi_j - \pi_m = f_{im} - f_{ij}$, which equals f_{im} at x . Same answer.

6 side3: a triangle crossing two bisectors

Signature: `side3(pi, wi, pj, wj, pk, wk, pm, wm, q0, q1, q2).`




Figure 2: An F vertex: where the plane of the triangle meets two bisectors at once. Drawn in the triangle's plane, with the numbers of the worked example (§6.3).

6.1 Derivation

Nine f values, three per corner q_r :

$$a_r = f_{ij}(q_r), \quad c_r = f_{ik}(q_r), \quad b_r = f_{im}(q_r), \quad r = 0, 1, 2.$$

The lemma with $k = 2$, $h_1 = f_{ij}$, $h_2 = f_{ik}$, $g = f_{im}$:

$$f_{im}(x) = \frac{\begin{vmatrix} b_0 & b_1 & b_2 \\ a_0 & a_1 & a_2 \\ c_0 & c_1 & c_2 \end{vmatrix}}{\begin{vmatrix} 1 & 1 & 1 \\ a_0 & a_1 & a_2 \\ c_0 & c_1 & c_2 \end{vmatrix}}. \quad (5)$$

$$\text{sign } f_{im}(x) = \text{sign}(N) \cdot \text{sign}(D), \quad N \text{ of degree } 6, \quad D \text{ of degree } 4.$$

6.2 Structure: share the bottom rows

Both determinants have the same bottom two rows. Expand both along the top row. The cofactors are the same three 2×2 minors of the a and c rows:

$$\mu_0 = a_1 c_2 - a_2 c_1, \quad \mu_1 = a_0 c_2 - a_2 c_0, \quad \mu_2 = a_0 c_1 - a_1 c_0,$$

and then

$$N = b_0 \mu_0 - b_1 \mu_1 + b_2 \mu_2, \quad D = \mu_0 - \mu_1 + \mu_2. \quad (6)$$

So the predicate costs nine f values, three minors (six products), and two cheap combinations. In the exact path this matters: the minors are the expensive products, and this way they are computed once instead of twice.

An alternative denominator. Subtract column 0 from columns 1 and 2 of D . The determinant is unchanged, and the row of ones becomes $(1, 0, 0)$:

$$D = (a_1 - a_0)(c_2 - c_0) - (a_2 - a_0)(c_1 - c_0).$$

As in **side2**, each difference is a dot product of coordinate differences, for example $a_r - a_0 = 2(q_r - q_0) \cdot (p_j - p_i)$. That is exactly the oracle's 2×2 system matrix, so D is its determinant. Same value, tighter error bound, optional.

6.3 Worked example

The same numbers as the hand-written test in `test_side3.cpp`.

$p_i = 0$, $p_j = (2, 0, 0)$, $p_k = (0, 2, 0)$, $p_m = (1, 1, 2)$, no weights, triangle $q_0 = (-4, -4, 0)$, $q_1 = (8, -4, 0)$, $q_2 = (-4, 8, 0)$ (figure 2). With $f_{ia}(q) = 2q \cdot p_a - \|p_a\|^2$:

$$\begin{aligned} f_{ij}(q) &= 4q_x - 4 : & a &= (-20, 28, -20), \\ f_{ik}(q) &= 4q_y - 4 : & c &= (-20, -20, 28), \\ f_{im}(q) &= 2(q_x + q_y + 2q_z) - 6 : & b &= (-22, 2, 2). \end{aligned}$$

The minors: $\mu_0 = 28 \cdot 28 - (-20)(-20) = 384$, $\mu_1 = (-20)(28) - (-20)(-20) = -960$, and $\mu_2 = (-20)(-20) - (28)(-20) = 960$. So

$$D = 384 + 960 + 960 = 2304, \quad N = (-22)(384) - 2(-960) + 2(960) = -4608,$$

and $f_{im}(x) = -2$: answer -1 . The alternative gives $D = 48 \cdot 48 - 0 \cdot 0 = 2304$.

Check: the vertex is $x = (1, 1, 0)$, with $\pi_i(x) = 2$ and $\pi_m(x) = 0 + 0 + 4 = 4$, a difference of -2 . ✓

6.4 Symmetries the tests check

- **Any order of the corners.** A permutation of the corners permutes the columns of both determinants in the same way, so both are multiplied by the same ± 1 .
- **Any of i, j, k called i .** The vertex is the point where all three power distances are equal, so $f_{im}(x) = f_{jm}(x) = f_{km}(x)$. The test rotates (i, j, k) and swaps j and k . This is a real check of the implementation, not of the formula: all nine f values change.

6.5 Cost of the exact path

This is the most expensive predicate: degree 6, with expansions of the f values multiplied three deep. The product of two `DynamicExpansions` compresses its result, which keeps each level short. Without that, every level would be slower than the one before. Only near-degenerate inputs reach this path. The filter decides the rest.

7 side4: a Voronoi vertex

Signature: `side4(pi, wi, pj, wj, pk, wk, pl, wl, pm, wm)`.

The vertex x has equal power distance to i, j, k and l . In three dimensions it is the centre of the power sphere through the four seeds. There is no domain feature, so the lemma does not apply. Instead the three bisector equations are solved directly, and the solution is substituted, again without ever being computed.

7.1 Work relative to p_i

Put the origin at p_i . Let $d_a = p_a - p_i$ and $x' = x - p_i$. Then

$$f_{ia}(x) = \|x'\|^2 - w_i - \|x' - d_a\|^2 + w_a = 2x' \cdot d_a - c_a, \quad \boxed{c_a = \|d_a\|^2 + w_i - w_a.}$$

Compare with (1): the large, cancelling $\|p_a\|^2 - \|p_i\|^2$ has become $\|d_a\|^2$, a square of differences. Everything below is built from the d_a and the c_a , so it is translation-invariant, for the same reasons as (3).

7.2 Derivation

The vertex satisfies $f_{ia}(x) = 0$ for $a = j, k, l$, that is

$$2d_a \cdot x' = c_a, \quad a = j, k, l.$$

Consider the 4×4 determinant whose rows are $(2d_a, c_a)$ for $a = j, k, l, m$. From the last column, subtract x'_1 times column 1, x'_2 times column 2 and x'_3 times column 3. This does not change the determinant. Each last entry becomes $c_a - 2d_a \cdot x' = -f_{ia}(x)$, which is 0 for $a = j, k, l$ and $-f_{im}(x)$ for $a = m$:

$$\begin{vmatrix} 2d_j & c_j \\ 2d_k & c_k \\ 2d_l & c_l \\ 2d_m & c_m \end{vmatrix} = \begin{vmatrix} 2d_j & 0 \\ 2d_k & 0 \\ 2d_l & 0 \\ 2d_m & -f_{im}(x) \end{vmatrix} = -f_{im}(x) \begin{vmatrix} 2d_j \\ 2d_k \\ 2d_l \end{vmatrix},$$

expanding along the last column, whose only non-zero entry sits at position (4,4) with cofactor sign $+$. Each d_a is a row of three numbers. Now drop the 2s. Factoring 2 out of each of the first three columns gives 2^3 on both sides, which cancels *exactly*, not just in sign:

$$f_{im}(x) = - \frac{\begin{vmatrix} d_j & c_j \\ d_k & c_k \\ d_l & c_l \\ d_m & c_m \end{vmatrix}}{\begin{vmatrix} d_j \\ d_k \\ d_l \end{vmatrix}}. \quad (7)$$

$$\text{sign } f_{im}(x) = - \text{sign}(N) \cdot \text{sign}(D), \quad N = \det_4 (\text{degree } 5), \quad D = \det_3 (\text{degree } 3).$$

Mind the minus.

Degrees: three columns of d entries (degree 1) and one column of c entries (degree 2) make N degree $1 + 1 + 1 + 2 = 5$. D is degree 3.

7.3 Structure: four triple products

Write $T(a, b, c) = \det[d_a; d_b; d_c] = d_a \cdot (d_b \times d_c)$, the triple product. Expand N along its last column. The cofactor signs for rows 1..4 of column 4 are $-, +, -, +$:

$$N = -c_j T(k, l, m) + c_k T(j, l, m) - c_l T(j, k, m) + c_m T(j, k, l), \quad D = T(j, k, l). \quad (8)$$

So the predicate is four c values and four triple products, one of which is D itself. Get one cofactor sign wrong and the answer is wrong only on *some* inputs. The hand-worked tests and the permutation test are what catch that.

7.4 The orient3d trap

T has the same shape as `orient3d_det`, so reusing it is tempting. But `orient3d` follows mesh-dev's convention, not this formula's:

$$\begin{aligned}\text{orient3d}(a, b, c, d) &= (a - d) \cdot ((c - d) \times (b - d)), \\ \text{so } \text{orient3d}(p_j, p_k, p_l, p_i) &= d_j \cdot (d_l \times d_k) = -T(j, k, l).\end{aligned}$$

Reusing it is fine if the sign is accounted for. Writing T directly is three lines and has no trap.

7.5 Worked example

The same numbers as the hand-written test in `test_side4.cpp`.

$p_i = 0$, $p_j = (2, 0, 0)$, $p_k = (0, 2, 0)$, $p_l = (0, 0, 2)$, $p_m = (3, 3, 3)$, no weights. Then $d_a = p_a$, $c_j = c_k = c_l = 4$ and $c_m = 27$. The triple products:

$$\begin{aligned}T(j, k, l) &= (2, 0, 0) \cdot ((0, 2, 0) \times (0, 0, 2)) = (2, 0, 0) \cdot (4, 0, 0) = 8, \\ T(k, l, m) &= (0, 2, 0) \cdot ((0, 0, 2) \times (3, 3, 3)) = (0, 2, 0) \cdot (-6, 6, 0) = 12, \\ T(j, l, m) &= (2, 0, 0) \cdot (-6, 6, 0) = -12, \\ T(j, k, m) &= (2, 0, 0) \cdot ((0, 2, 0) \times (3, 3, 3)) = (2, 0, 0) \cdot (6, 0, -6) = 12.\end{aligned}$$

By (8), $N = -4 \cdot 12 + 4 \cdot (-12) - 4 \cdot 12 + 27 \cdot 8 = 72$, $D = 8$, and $f_{im}(x) = -72/8 = -9$: answer -1 .

Check: the vertex is $x = (1, 1, 1)$, with $\pi_i(x) = 3$ and $\pi_m(x) = 4 + 4 + 4 = 12$, a difference of -9 . ✓ The same test then gives p_m a weight of 10. Only c_m changes, to $27 - 10 = 17$, so $N = -144 + 136 = -8$ and $f_{im}(x) = +1$: the weight pulls the vertex to m 's side.

7.6 Symmetries the tests check

All 24 orders of (i, j, k, l) must give the same sign, since the vertex and its power distance are the same. Permuting j, k, l among themselves permutes rows of both determinants together, so it is easy. Exchanging i with one of them changes *every* d_a and c_a . That this still gives the same sign is true of the mathematics, and it is the strongest check there is of the sign bookkeeping in (7) and (8).

8 Evaluating them: fast when easy, exact when not

Every predicate has the same skeleton as `orient2d`, `orient3d` and `side1`.

8.1 One formula, several number types

Write each formula (the f values, the minors, the triple products, N and D) once, as a template over a number type T . The same text then means three things:

T	what the formula computes
<code>double</code>	a rounded value: fast, and wrong exactly where it matters
<code>Bounded</code>	a rounded value <i>and</i> a certified bound on its error: the filter
<code>DynamicExpansion</code>	the exact value, as a list of non-overlapping doubles: the exact path

8.2 The filter

A **Bounded** carries a value v and an error e with $|v - \text{exact}| \leq e$. Inputs have $e = 0$. With $u = 2^{-53}$, each operation propagates the bound:

$$a \pm b : e = e_a + e_b + u|v|, \quad a \cdot b : e = |a|e_b + |b|e_a + e_a e_b + u|v|, \quad -a : e = e_a,$$

inflated by $1 + 2^{-48}$ to cover the rounding of the bound itself. At the end, **certain_sign** says:

- $|v| > e$: the sign of v is the true sign;
- $v = 0$ and $e = 0$: the value is exactly zero;
- otherwise: don't know.

The filter can be unsure. It cannot be wrong. Because the bound is built by the operators, it is sound for any formula, including these new ones, with no hand derivation of error constants.

8.3 The decision

1. Evaluate N and D with **Bounded**.
2. If **certain_sign** says N is certainly 0, the answer is 0 whatever D is: $D \neq 0$ is the precondition, so D 's sign is not needed. Count one **filtered** call. (This is what the code does; it does not certify D first.)
3. Otherwise, if **certain_sign** decides *both*, the answer is the product of the two signs. Count one **filtered** call.
4. Otherwise, evaluate N and D exactly with **DynamicExpansion**, take the sign of each (the sign of the largest component), and return their product. Count one **exact** call.

The side4 minus. In this document **side4**'s N is $\det_4$ itself, so its answer is $-\text{sign } N \cdot \text{sign } D$. The code folds the minus into the numerator instead: **side4_terms** returns $\{-\det_4, \det_3\}$, and the decision above, with no extra minus, is then the same for all four predicates.

Two refinements are possible later, and neither is needed for correctness. If the filter certified D but not N , only N needs the exact path. And the f values could be computed once and shared between N and D , which the minor structure of **side3** already does.

Why both signs. It is tempting to certify only N and assume $D > 0$. As shown in §3, the sign of D depends on the order of the corners (or, for **side4**, on the orientation of p_i, p_j, p_k, p_l), which is arbitrary. So it is part of the answer.

The precondition. $D \neq 0$. If the exact D is zero, the caller asked about a vertex its equations do not define uniquely. That is a bug in the caller, and an assertion should say so. The tests keep assertions on in every build type.

8.4 Why DynamicExpansion

Expansion<N> carries a worst-case capacity in its type. A sum of M and N components needs $M + N$, and a product needs $2MN$. That works for **orient3d** (192 components, 1.5 KB), but the bounds multiply with the degree. **side2** would need about ten thousand components, and **side3** millions, far more than a stack holds. Those are worst cases. The real length of an expansion depends on how many bits the value actually has, which is usually far fewer, and known only once the value is computed. **DynamicExpansion** keeps its first 32 components inline and moves to the heap beyond that. An exact evaluation therefore avoids the heap only if *every intermediate value* stays within 32 components, not merely the final result: products and the merge in a sum can be longer than what is left at the end.

8.5 Counters

Each predicate has its own thread-local `PredicateCounts`: `side2_counts()`, `side3_counts()` and `side4_counts()`, like `side1_counts()`. Every call increments exactly one of `filtered` and `exact`. The tests check that the two add up to the number of calls, and they use the ratio to check that the filter does its job.

9 How they are tested

9.1 The oracle

The oracle (`tests/side_support.h`) shares no formula with the predicates. It builds the vertex’s defining linear system (the bisector equations, in barycentric or Cartesian unknowns), solves it exactly in GMP rationals by Gaussian elimination, and compares $\pi_i(x)$ with $\pi_m(x)$ straight from the definition $\|x - p\|^2 - w$. It returns “no vertex” when the system is singular, and the tests then skip the case. Two derivations from different starting points that agree on hundreds of thousands of inputs is the evidence.

9.2 The input families

family	what it is for
random	coordinates with random signs, full significands and exponents in $[-20, 20]$, half the seeds weighted. Cancellation at every scale. The filter should decide nearly all of them
lattice	small integers in $[-2, 2]$. Ties are common, and the answer must then be exactly 0
far lattice	the same, moved by 2^{20} and by $10^6 + 0.5$. Same answers, ties included, under heavy cancellation
nudged ties	a lattice tie with one non-zero coordinate moved by one ulp. Tiny non-zero answers. They must mostly reach the exact path, and constructing the vertex in doubles must get some of them wrong (otherwise the test would be too easy)

9.3 Thresholds

The required values are the suites’ thresholds. They were set on 1 October 2026 below what a throwaway reference implementation reached (with a different random seed, whose results were not kept). The “measured” values are those of the actual `side.h`, on 5 October 2026, on the suites’ own input families and seeds, from `build/rvd-measure-side` (`tests/measure_side.cpp`; setup and timings in `documents/arithmetic`, section 9). They are deterministic, so running the program again reproduces them exactly.

	filter on random (required / measured)	lattice ties (measured)	nudged reaching exact (required / measured)	doubles wrong (measured)
side1	99% / 100.00%	4.54%	80% / 88.1%	76.3%
side2	99% / 100.00%	3.26%	80% / 91.6%	71.8%
side3	95% / 98.19%	2.77%	80% / 91.6%	67.7%
side4	99% / 99.99%	3.98%	80% / 93.6%	65.2%

Not every nudged tie needs the exact path. A nudge can move a value far enough from zero, relative to its error bound, for the filter to certify it. “Doubles wrong” is the share of nudged ties

on which constructing the vertex in doubles, then comparing the two power distances, gets the sign wrong; the suites only require it to be above zero.

10 A preview: ties

For now an exact tie returns 0, and the lattice tests expect that. A later step replaces 0 with *symbolic perturbation* (DESIGN.tex, section 6.4). The idea, so that it is not a surprise: imagine every weight increased by an infinitesimal ε_a , with $\varepsilon_0 \gg \varepsilon_1 \gg \varepsilon_2 \gg \dots$. Increasing w_a by ε_a lowers π_a by ε_a , so it shifts

$$f_{ia} \mapsto f_{ia} + (\varepsilon_a - \varepsilon_i).$$

Every f value moves by a constant, so the perturbed N is the true N plus terms in the ε s, whose coefficients are smaller determinants of the same family, with a row replaced by ones. When the true value is 0, the sign is that of the first non-zero coefficient, in order of priority. For **side1** this is simple: $f_{im}(q) = 0$ becomes $\varepsilon_m - \varepsilon_i$, whose sign is decided by which index is smaller. For the others the chain always ends at the vertex's own D , which is non-zero by the precondition. The details are worked out in `documents/arithmetic/arithmetic.tex`, section 8. Nothing in this document's formulas changes, but the formulas will be reused for the coefficients, which is one more reason to write them cleanly.

11 Checklist

- f values in the difference form (3), through `bisector_value`.
- The answer is $\text{sign } N \cdot \text{sign } D$, never $\text{sign } N$ alone. **side4** has an extra minus.
- **side2**: $N = b_0a_1 - b_1a_0$, $D = a_1 - a_0$. Both or neither negated.
- **side3**: three minors μ_r , shared by N and D ; cofactor signs $+, -, +$.
- **side4**: coordinates relative to p_i ; $c_a = \|d_a\|^2 + w_i - w_a$; cofactor signs $-, +, -, +$ on the column of cs ; **orient3d** is $-T$.
- Filter: an answer needs both N and D certain, or N certainly zero (then the answer is 0, by the precondition $D \neq 0$).
- One counter increment per call, in the right counter.
- Assert $D \neq 0$ on the exact path.
- A suite joins the build as soon as `sideK(` appears in `side.h`. From then on it has to compile and pass.
- Run one suite with `cmake -build build && build/rvd-tests-side2`.

Summary

	numerator N	denominator D	deg N	deg D
side1	$f_{im}(q)$	1	2	0
side2	$b_0a_1 - b_1a_0$	$a_1 - a_0$	4	2
side3	$b_0\mu_0 - b_1\mu_1 + b_2\mu_2$	$\mu_0 - \mu_1 + \mu_2$	6	4
side4	$\det_4[(d_a, c_a)]_{a=j,k,l,m}$, answer negated	$T(j, k, l)$	5	3